Catalog Reconciliation — Human TODO

Created: 2026-04-27 **Status:** Sync ran against `data_tables/real/366132845011`. Catalog now has 12 YAMLs. 24 columns + 1 table need descriptions. 5 profile-only tables + 26 profile-only columns need a keep/drop/mark decision. 2 aliases need a sanity-check.

Context — what sync already did

- Read all real CSVs under `data_tables/real/366132845011/` (skipped 1 `.xlsx`).
- Folded every real column it could match into an existing canonical YAML (added 36+ aliases across `bureau`, `cross_bu`, `payments`, `spends`, `model_scores`).
- For real tables that didn't substring-match any canonical: created `modelling_data.yaml` (the only one — all others folded into existing canonicals).
- Wrote new columns under their resolved canonical with `description_pending: true` and a regex-pattern draft (no LLM was available — `OPENAI_API_KEY` not set).

What's left is everything that needed human judgment.

🔴 P1 — Write descriptions (24 columns + 1 table)

Table description (1)

- ☐ **`config/data_profiles/modelling_data.yaml`** → set `description:` field. Columns are `cbr_score_max`, `exp_pif_max`, `ons_30_trd_mean` — looks like a model-output table.

Group A — `score_drivers` ranked driver names (20 cols, single pattern)

File: `config/data_profiles/score_drivers.yaml`. These look like top/bottom 5 contributors for two model scores (CDSS and TSR). One template description should work for all 20:

> *"Name of the rank-{N} {top|bottom} contributing variable to the {CDSS|TSR} model score for this case."*

Fill in `description:` and set `description_pending: false` for:

- ☐ `top_cdss1` `top_cdss2` `top_cdss3` `top_cdss4` `top_cdss5`
- ☐ `bottom_cdss1` `bottom_cdss2` `bottom_cdss3` `bottom_cdss4` `bottom_cdss5`
- ☐ `top_tsr1` `top_tsr2` `top_tsr3` `top_tsr4` `top_tsr5`
- ☐ `bottom_tsr1` `bottom_tsr2` `bottom_tsr3` `bottom_tsr4` `bottom_tsr5`

Group B — Replace weak regex drafts (3 cols)

- ☐ **`spends.Date`** — currently `"date value — verify format and semantics ..."`. Replace with the actual semantic (e.g., *"Posting date of the spend transaction (YYYY-MM-DD)."*).
- ☐ **`wcc_flags.Date`** — same. Likely *"Date the WCC flag was raised."*
- ☐ **`modelling_data.cbr_score_max`** — currently `"score value (agent-drafted, unverified)"`. Find the actual meaning (likely Consumer Bureau Risk score, max over a window).

Group C — Genuinely empty (2 cols)

- ☐ `modelling_data.exp_pif_max` — string. Meaning unclear from name. Likely Experian Paid-In-Full max.
- ☐ `modelling_data.ons_30_trd_mean` — int. Likely "onus 30-day delinquent tradelines, mean" — verify.

Group D — Empty col description, but col exists in real (1 col)

- ☐ `wcc_flags.Note` — free-text note column. *"Free-text note attached to the WCC flag."*

🟡 P2 — Verify two suspect aliases

These got auto-accepted at top during sync but their normalized names don't perfectly overlap with the canonical column. Both look correct on inspection — just need a yes/no:

- ☐ `bureau.avg_external_utilization` ← `Average External Utlization` — typo in real ("Utlization"); semantically the same as canonical `utilization`. **OK?**
- ☐ `cross_bu.past_delinquencies_12m` ← `Past Delinquencies (Last 12M)` — parenthetical phrasing differs but semantically the same. **OK?**

If either is wrong, edit the alias list manually in the YAML.

🟠 P3 — Keep / drop / mark profile-only entries

These exist in canonical but no real CSV has them. For each, decide:

- **Keep** — future real-data feed expected; leave as-is.
- **Drop** — outdated; delete from YAML.
- **Mark simulator-only** — keep for the data generator but tell the agent path to ignore. (No mechanism for this today — would need a `simulator_only: true` field + filter in `to_prompt_context`.)

5 tables (whole-table decisions)

- ☐ `cust_tenure`
- ☐ `income_dti`
- ☐ `model_scores` *(NB: `modelling_data` exists in real; possibly should be merged into this canonical instead of standing alone — see Open Question 1 below.)*
- ☐ `txn_monthly`
- ☐ `xbu_summary`

26 columns (in tables that have real data, but these specific cols don't)

`bureau` **per-tradeline columns** — real data is monthly aggregate, not per-tradeline. These 10 cols are simulator-only by design:

- ☐ `score_date`, `trade_type`, `balance`, `credit_limit`, `utilization`, `dpd_status`, `delinquency_amount`, `is_revolving`, `open_date`, `last_reported_date`

`cross_bu` (3 cols, possibly in `crossbu_merchants_data` under different names):

- ☐ `merchant_name`, `merchant_industry`, `merchant_charge_volume`

`score_drivers` (7 cols — older driver schema vs the new top/bottom cdss/tsr shape):

- ☐ `model_id`, `driver_rank`, `driver_variable`, `driver_direction`, `driver_value`, `driver_contribution`, `driver_description`

`spends` (3 cols):

- ☐ `spend_date`, `month`, `customer_industry`

`wcc_flags` (3 cols):

- ☐ `flag_type`, `trigger_date`, `severity`

How to execute

Path 1 — Interactive sync with LLM (recommended for P1)

```
export OPENAI_API_KEY=sk-...
python -m datalayer.sync
```

Walks you through every `description_pending: true` column with a fresh LLM draft. ENTER = accept. `e` = edit. `r` = regenerate. About 24 prompts.

Path 2 — Manual edits

Open each YAML in the IDE, fill in `description:`, flip `description_pending: false`. Best for the 20 `score_drivers` cols where the same template applies.

Path 3 — Hybrid `verify` CLI (not yet built)

A small subcommand that finds pending columns and prompts for text only (no LLM needed). Ask Claude to build it if you want this — ~50 lines.

For P2

Open `bureau.yaml` and `cross_bu.yaml`, find the two columns, decide. No code needed.

For P3

Decide your overall stance, then I can apply it as one batch edit:

- **"Drop all simulator-only"** → I'll remove the 5 tables + 26 cols across YAMLs in one pass.
- **"Keep all"** → no action.
- **"Mark with `simulator_only: true`"** → adds the field + needs a follow-up code change to filter them out of agent-facing catalog views.

Open Questions (worth deciding before merging anything)

1. **`model_scores` (canonical) vs `modelling_data` (real) — should they merge?** They didn't auto-fold because `modelscores` and `modellingdata` don't substring-match. But they're plausibly the same concept under different names. If yes → rename one and re-run sync, or hand-merge.

2. **`crossbu_cards_data` vs `crossbu_merchants_data` — both fold into `cross_bu`. Is that right?** Real has two separate files; canonical has one. Today their columns all alias under `cross_bu`. Acceptable, but if they're meaningfully distinct in the source system, maybe split into `cross_bu_cards` and `cross_bu_merchants` in canonical.

3. **`payments_success` + `payments_returns` → both fold into `payments`.** Same question.

4. **`demographics_data.xlsx` is invisible** — gateway is CSV-only. Either convert once to CSV or extend gateway. Currently no demographics info reaches the catalog at all.

When you're done

After P1+P2 are clean, verify:

```
python -c "
import yaml
from pathlib import Path
pending = []
for f in Path('config/data_profiles').glob('*.yaml'):
    p = yaml.safe_load(open(f))
    for col, spec in (p.get('columns') or {}).items():
        if isinstance(spec, dict) and spec.get('description_pending') is
True:
            pending.append(f'{p[\"table\"]}.{col}')
print(f'{len(pending)} columns still pending:')
for c in pending: print(' -', c)
"
```

Target: zero pending.

Then commit:

```
git add config/data_profiles/
git commit -m "catalog: verify descriptions for real-data columns from case
366132845011"
```