AgenticSys v2 — OpenAI Agents SDK Wiring Boundary

What the SDK auto-wires vs what we wire ourselves

AgenticSys v2 — post openai-agents migration

2026-04-28

Contents

Overview 1

Part 1 — What the SDK wires automatically 3

Part 2 — What the SDK does NOT wire — we do, ourselves 3

2.1 The LLM client itself (`llm/firewall_client.py`) 4
2.2 Inter-agent transit redaction (`case_agents/redacting_tool.py`) 4
2.3 Beta trace-extraction fallback (`orchestrator/orchestrator.py`) 4
2.4 Tool dependency injection for data tools (`tools/data_tools.py`) 4
2.5 Skill / instructions composition 4
2.6 The pre-orchestrator pipeline (`main.py`) 4
2.7 ChatAgent (entirely outside the SDK) 5
2.8 EventLogger 5
2.9 The Orchestrator class wrapper 5

Mental model 5

Quick reference — where to look 6

Overview

This document draws a clean line between two layers of AgenticSys v2:

- **What the OpenAI Agents SDK does for us automatically** — the agent loop, tool dispatch, parallelism, structured-output parsing.
- **What the SDK does NOT do — and how we wire it ourselves** — the LLM client layer, inter-agent transit redaction, beta fallback, dependency injection, instructions composition, the pre/post pipeline, ChatAgent, and the Orchestrator wrapper class.

The goal is to make it obvious where to look when something needs tuning or debugging.

> The diagram above is the visual companion to the textual sections below. Blue boxes are our code; green is SDK-managed (auto-wired); amber is the firewall client we wire on the LLM transport layer; rose is the external OpenAI API. The dashed green container marks the SDK boundary — everything inside `Runner.run` is the SDK's responsibility; everything outside is ours. Regenerate with `python brainstorm/_sdk_wiring_diagram.py`.

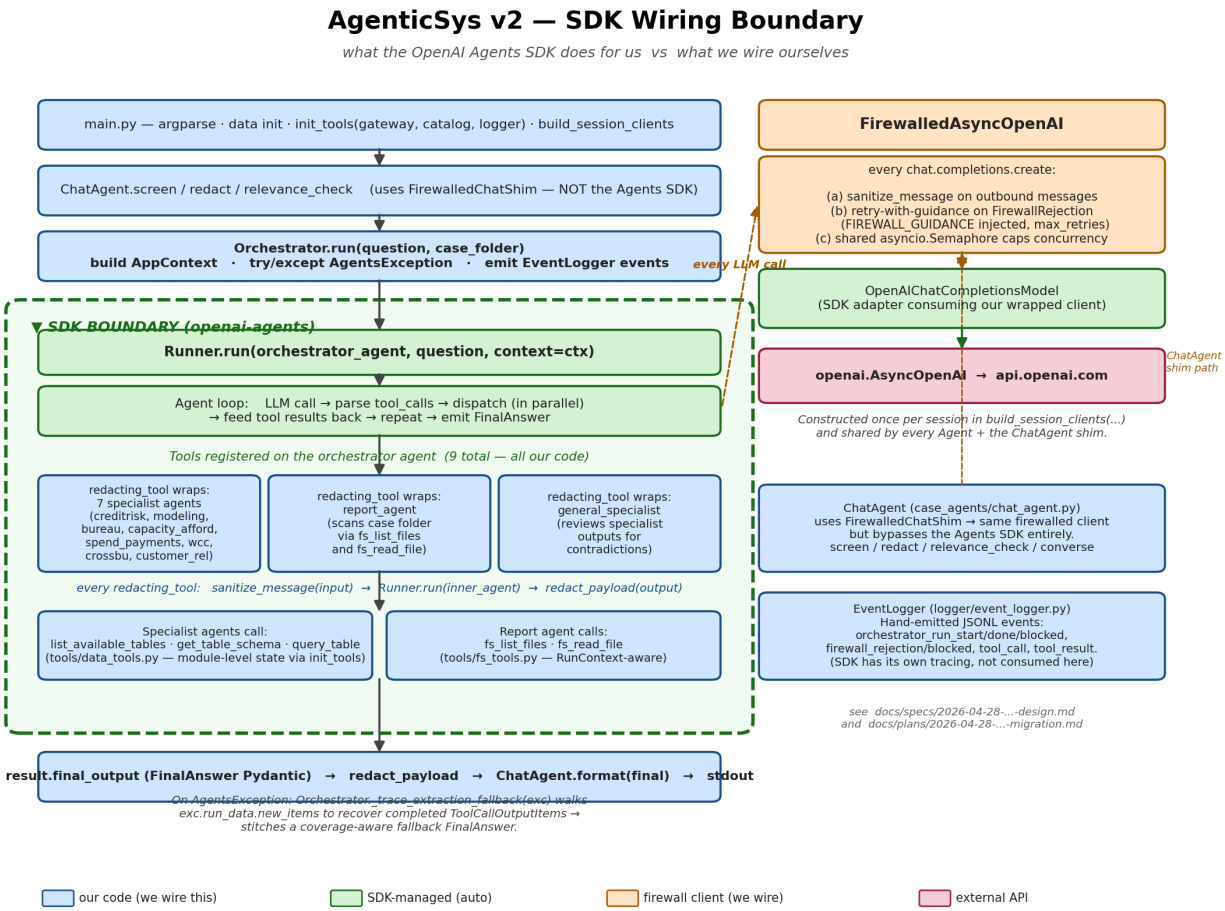

Figure 1: SDK wiring boundary diagram

Part 1 — What the SDK wires automatically

The SDK runs the **agent loop** for us — that is its core value. Inside one `Runner.run(agent, input, context=...)` call, the SDK does all of the following without our code touching it:

What the SDK does	Concretely
The loop itself	LLM call -> check for tool_calls -> dispatch tools -> feed results back as messages -> repeat until LLM emits a final structured output (or hits `max_turns`).
Tool dispatch	Given an LLM-emitted `tool_call`, finds the matching `FunctionTool` by `name`, deserializes JSON args against `params_json_schema`, calls `tool.on_invoke_tool(ctx, args)`, serializes the return into the next LLM message.
Parallel tool execution	When the LLM emits multiple `tool_calls` in one response, the SDK schedules them concurrently (`asyncio.gather` under the hood). No `await asyncio.gather` in our code — the SDK does it.
`ctx` injection	The `context=AppContext(...)` argument to `Runner.run` is auto-threaded into every tool whose first parameter is typed `ctx: RunContextWrapper[AppContext]`. The `ctx` param is also stripped from the JSON schema the LLM sees.
Structured output parsing	`output_type=AgentOutputSchema(FinalAnswer, ...)` tells the SDK to (a) pass `FinalAnswer`'s JSON schema as OpenAI's `response_format`, (b) parse the response JSON into a `FinalAnswer` Pydantic instance, (c) expose it as `result.final_output`.
Agent-as-tool	`agent.as_tool(tool_name, tool_description)` produces a `FunctionTool` that, when invoked, internally runs `Runner.run(inner_agent, ...)`. This is how nested agent dispatch works.
Run-history bookkeeping	The SDK accumulates `RunItems` (`ToolCallItem`, `ToolCallOutputItem`, `MessageOutputItem`) as turns progress, and exposes them on `result.new_items`. On exception, it stashes them under `exc.run_data.new_items` for partial-trace inspection.
Model invocation	`OpenAIChatCompletionsModel(model="gpt-4o", openai_client=...)` is the SDK-provided adapter that converts the SDK's internal message format into the OpenAI `chat.completions.create` API call.

Part 2 — What the SDK does NOT wire — we do, ourselves

Anything outside the agent loop, anything customizing how the LLM client behaves, and anything that touches infrastructure beyond the SDK's awareness.

2.1 The LLM client itself (`llm/firewall_client.py`)

The SDK calls `openai_client.chat.completions.create(...)` to talk to OpenAI. We *substitute* a wrapped client (`FirewalledAsyncOpenAI`) that:

- Redacts PII from every outbound message before forwarding to the real client
- Catches `FirewallRejection` and retries with `FIREWALL_GUIDANCE` injected into the system message (up to `max_retries`)
- Acquires a shared `asyncio.Semaphore` around each request so global in-flight calls are capped

The SDK has no idea this layer exists. It just sees an `AsyncOpenAI`-shaped object.

2.2 Inter-agent transit redaction (`case_agents/redacting_tool.py`)

When the orchestrator's LLM calls a specialist tool, the SDK passes the LLM-generated string straight through to the inner agent. The SDK does not transform or sanitize. We add that boundary ourselves: `redacting_tool` is a `@function_tool` that wraps `Runner.run(inner_agent, ...)` with `sanitize_message` on input and `redact_payload` on output. This is why every sub-agent in the orchestrator's tool list is wrapped — without `redacting_tool`, the inter-agent edges would be unredacted.

2.3 Beta trace-extraction fallback (`orchestrator/orchestrator.py`)

The SDK provides the data (`exc.run_data.new_items` after `AgentsException`), but the *strategy* — "if the orchestrator fails after some specialists ran, recover their `SpecialistOutputs` and stitch a coverage-aware fallback `FinalAnswer`" — is hand-written in `Orchestrator._trace_extraction_fallback`. The SDK gives us the parts; we decide what to do with them.

2.4 Tool dependency injection for data tools (`tools/data_tools.py`)

The SDK supports per-request `RunContext` (which we use for `fs_tools`), but for `data_tools` we kept the legacy module-level state pattern:

- `_gateway`, `_catalog`, `_logger` are module-level globals
- `main.py` calls `init_tools(gateway, catalog, logger)` once at session start before any `Runner.run`
- The decorated `query_table` etc. read from those globals at call time

The SDK does not know about this — to it, `query_table` is just a function returning a string. The dependency injection is happening in our module's lexical scope, invisible to the runtime.

2.5 Skill / instructions composition

Every Agent's `instructions=` is a static string we build at module-import time by concatenating markdown bodies (`team_construction.md` + `data_catalog.md` + ...). The SDK never reads files, never composes, never refreshes — it just passes that exact string as the system prompt on every turn. If you edit a markdown file, you must re-import the module to see the change.

The composition functions (`_compose_orchestrator_instructions`, `_compose_instructions` in `specialist_agent.py`, etc.) live entirely outside the SDK's awareness.

2.6 The pre-orchestrator pipeline (`main.py`)

Everything that happens before `await Runner.run(...)` is hand-wired:

- CLI argparse
- Data layer init: `LocalDataGateway`, `DataCatalog`, `DataGenerator`

- Pillar config loading from YAML
- `init_tools(...)` to set the data-tool module globals
- `ChatAgent.screen(...)` for question-time content screening (uses `FirewalledChatShim`, **not** the Agents SDK)
- Case-folder resolution

After `Runner.run(...)` returns, we also do:

- `redact_payload(result.final_output)` (one last redaction at the SDK -> caller boundary)
- `ChatAgent.format(final)` to render the `FinalAnswer` as user-facing markdown

None of these run inside the SDK loop.

2.7 ChatAgent (entirely outside the SDK)

`case_agents/chat_agent.py`'s `screen`, `redact`, `relevance_check`, `converse` methods do not use `Agent`/`Runner` at all. They make plain `chat.completions.create` calls via the `FirewalledChatShim` (`llm/factory.py`). The shim mimics the legacy `FirewalledModel.ainvoke(system_prompt, user_message)` interface so ChatAgent did not need to be rewritten. From the SDK's perspective, ChatAgent is invisible — it is a separate code path that happens to share the firewalled `AsyncOpenAI` client.

2.8 EventLogger

Our JSONL logger (`logger/event_logger.py`) is hand-written and called explicitly from `Orchestrator.run` (`orchestrator_run_start`, `orchestrator_run_done`, `orchestrator_run_blocked`), `FirewalledAsyncOpenAI` (`firewall_rejection`, `firewall_blocked`), and `data_tools` (`tool_call`, `tool_result`). The SDK has its own tracing system (which we do not currently consume); both run side-by-side.

2.9 The Orchestrator class wrapper

The SDK gives you `Runner.run(agent, input)`. That is it. We wrap it in `Orchestrator.run(question, case_folder, report_agent=None)` to:

- Build `AppContext` from session state
- Catch `AgentsException` for the beta fallback
- Apply the final redaction
- Emit our `EventLogger` events

The SDK has no concept of an "Orchestrator class" — that is our framing on top.

Mental model

```
+------------------------------ our code --------------------------------+
|  main.py: argparse, data init, init_tools, ChatAgent screening       |
|     |                                                                 |
|     v                                                                 |
|  Orchestrator.run                                                     |
|     |  builds AppContext, catches AgentsException                     |
|     v                                                                 |
|  +---------------- SDK boundary -----------------+                   |
|  |  Runner.run(orchestrator_agent, ...)          |                   |
|  |     |                                          |                    |
|  |     v                                          |                    |
|  |  agent loop:                                   |                    |
```

```
|  |    LLM call ----> openai_client.chat...create  |                    |
|  |    tool_call dispatch                          |                    |
|  |    parallel scheduling                         |                    |
|  |    structured output parsing                   |                    |
|  |     |                                          |                    |
|  |     v                                          |                    |
|  |  result.final_output: FinalAnswer              |                    |
|  +------------------------------------------------+                    |
|     |                                                                 |
|     v                                                                 |
|  redact_payload, ChatAgent.format, return to caller                   |
+-----------------------------------------------------------------------+

Side channel running THE WHOLE TIME:
  openai_client = FirewalledAsyncOpenAI(real_AsyncOpenAI, firewall)
  |- redacts every outbound message
  |- retries on FirewallRejection with FIREWALL_GUIDANCE
  +- shared semaphore caps concurrency

Each tool the SDK dispatches:
  - data_tools: read module-level globals set by init_tools
  - fs_tools: read AppContext via RunContextWrapper
  - sub-agents (via redacting_tool): sanitize input, run inner agent,
    redact output before SDK feeds it back to outer LLM
```

The SDK is the **dispatcher** — it owns the loop, parallelism, JSON marshaling, and tool resolution. Everything else (transport-layer behavior, inter-agent boundary security, instruction composition, dependency injection patterns, the upstream / downstream pipeline, observability) is ours.

Quick reference — where to look

If you need to tune…	Look in
Orchestrator's high-level prompt	`case_agents/orchestrator_agent.py` -> `_compose_orchestrator_instructions` (incl. inline `TOOL-USE DISCIPLINE` block)
The skill prose itself	`skills/workflow/*.md`, `skills/domain/*.md`
Specialist instruction template	`case_agents/specialist_agent.py` -> `_compose_instructions`
What each tool looks like to the LLM	Tool docstring (`tools/data_tools.py`, `tools/fs_tools.py`) and tool descriptions in the orchestrator factory
HTTP-layer policy (redaction, retry, throttling)	`llm/firewall_client.py`
Inter-agent redaction	`case_agents/redacting_tool.py`
Beta fallback strategy	`orchestrator/orchestrator.py` -> `_trace_extraction_fallback`
Data tool dependencies (gateway/catalog/logger)	`tools/data_tools.py` -> `init_tools`, called once from `main.py`
Per-request context (case_folder, gateway)	`case_agents/app_context.py` -> `AppContext`, threaded by `Runner.run(context=...)`
Pre/post-orchestrator pipeline	`main.py` (`amain`, `_screen_and_run`, `run_question`)
ChatAgent's LLM calls	`llm/factory.py` -> `FirewalledChatShim`

If you need to tune…	Look in
Model + decoding choice	`llm/factory.py` -> `build_session_clients(model_name=...)` and `main.py` `--model` flag