

Timeouts and retries

2026-08-16. Numbers are the shipped [config/tuning.yaml](#).

Four independent layers can each abandon work, and each has its own retry. They nest, and their retries **multiply** — which is the part that is easy to get wrong when tuning any single number.

****Three places now treat a slow call as *suspect* rather than as work in progress**, and they are the same mechanism at different heights: abandon the attempt, re-issue once, never loop.**

layer	fence	then	backends	section
CALL	40s	re-issue on a 60s budget	safechain	"The call layer"
SCREEN phase	10s	re-issue on what is left of the phase's 30s	both	"Why screen needs its own fence"
ROUND-1 watchdog	25s	replan -- and the deadline is DISARMED on the last attempt	both	"The round-1 watchdog"

The lower two exist because the CALL layer cannot reach them: its fence is 40s and both phases are budgeted under that, so the phase fence always fires first.

(Diagrams are plain ASCII on purpose: box-drawing glyphs are outside Courier's WinAnsi encoding and render substituted or blank in the PDF.)

The layers

```
+ = TURN ===== turn_wall_clock_s = 360s =+
| asyncio.wait_for on the whole turn. On expiry the coroutine is CANCELLED, |
| so every agent still in flight dies with it -- they report `cancelled`, |
| not a timeout of their own. no retry: the turn ends |
|
| +- PHASE ----- one fence per kind of agent -----+
| | screen 30s (stall fence 10s, INSIDE the 30) |
| | orch_plan 25s (the round-1 watchdog -- it REPLANS) |
| | reviewer 120s specialist 240s |
| | distiller 120s distiller_drain 60s report_agent 150s |
| |
| | Expiry = that agent failed. The turn continues, degraded. |
| | Two retry here: orch_plan (replan) and screen (re-issue). `reviewer` |
| | does NOT -- which is why its fence is wide enough to outlast a stall |
| | and let the CALL layer below do the retrying instead. |
| |
| | +- AGENT ----- _MAX_SPECIALIST_ATTEMPTS = 2 -----+
| | | Re-runs the whole agentic loop when the ANSWER is wrong: |
| | | ungrounded_retry built on a failed tool call |
| | | absence_reread denies what the tools returned |
| | | max_turns_retry blew its turn budget |
| | |
| | | +- CALL ----- one HTTP request to the model -----+
| | | | stall fence 40s -> abandon, re-issue ONCE |
| | | | full budget 60s -> TimeoutError + safechain_retry_ |
| | | | stalled (the falsifying datum) |
| | | | HTTP 401 -> refresh token, retry IN PLACE |
| | | | 403 / 400 -> FirewallRejection, escalates to the |
| | | | guidance loop below |
| | |
| | | +-----+
| | | +-----+
| | +-----+
| +-----+
+=====
```

```

+- REJECTION (wraps the call) ----- firewall.max_retries = 2 -----+
| 403 compliance block / 400 bad request are DETERMINISTIC, so an      |
| identical retry is guaranteed to fail the same way. This loop        |
| retries with a CHANGED input instead -- `_inject_guidance` appends   |
| firewall guidance telling the model what to remove. 401 is the        |
| opposite case (a stale token, transient) and is retried in place.     |
+-----+

```

Also in the path, but not a fence: the firewall semaphores (specialist 8, orchestrator 4).

The call layer, in detail

Added 2026-08-13, for safechain only. The screen phase carries the same idea at a different layer (see below); everything about the reasoning is shared.

```

issue request
|
+-- returns < 40s -----> ANSWER  (p99 is 7.5s)
|
+-- still running at 40s
    | the call is not slow, it is WEDGED
    | log: safechain_call_stalled
    v
  abandon it  (ainvoke is genuinely cancellable, so the
               request aborts -- no orphaned thread)
    |
  issue a FRESH request, full 60s budget
    |
    +-- returns -----> ANSWER  (~45s total)
    |
    +-- stalls too -----> FAILS at 100s
        | logs safechain_retry_stalled
        +-- exactly one extra attempt; never a loop

```

The budget bets that the retry ESCAPES. The alternative was a budget above the 126-131s plateau so a stalled retry could ride it out — safer, but it makes every persistent wedge cost ~170s. The bet is supported by concurrent divergence: `distiller.bureau` returned in 5.8s while `distiller.modeling` hung 120s in the same pool at the same moment, so the wedge belongs to the REQUEST and a fresh one should not inherit it. Given that, the budget only has to cover a healthy call — p99 7.5s in dev, 2.1-5.6s in prod, so 60s leaves ~8x margin.

	before	after
retry escapes the stall	~130s	~45s
retry stalls too	~130s	fails at 100s, logged
call is healthy	unchanged	unchanged, issued once

The bet is falsifiable, and recorded in two places because they have different readers. The session JSONL carries `safechain_call_stalled` and `safechain_retry_stalled` for a human reading one run. The node trace carries the same thing as tags, because AgenticEval scores from the trace DB and never reads the JSONL — a stall visible only in the log would be invisible to every scored run:

```

node      specialist.modeling.round_1
tags      ["stall_retry"]           call abandoned and re-issued
          ["stall_retry_failed"]   the re-issue ALSO timed out
extra_json {"stall_retry_after_s": 40, "retry_budget_s": 60}

```

Tagged on the ROUND, so a stall attributes to a specific agent. Zero `stall_retry_failed` means the budget can drop further; anything non-trivial means the wedge survives re-issue and the budget must go back above the plateau.

Why 40s and not tighter. Per-call latency is p99 7.5s over 88 dev calls and 2.1-5.6s in prod, so 40 is ~5x the observed worst case: it cannot abandon a healthy call, and it only governs how long a STALL burns before the retry.

`SAFECHAIN_STALL_RETRY_S=0` disables it. The first attempt is clamped to `min(stall_fence, full_budget)`, so lowering `SAFECHAIN_CALL_TIMEOUT_S` really does lower it — unclamped, the fence silently outlived a tighter setting.

Why this layer exists

Measured in the private env: safechain calls do not run slow, they **stall**. In one turn `distiller.bureau` returned in 5.8s while `distiller.modeling` — concurrent, same pool, same model — hung. Every stall allowed to finish landed in a narrow band:

```
orchestrator.round_2 125.98s ok    spend_payments 129.50s ok
spend_payments       129.83s ok    crossbu        130.67s ok
```

while every failure died at **its own phase fence**, not at a natural duration:

```
report_agent          100.01s  <- report_agent_s was 100, now 150
distiller.modeling     120.01s  <- distiller_s = 120
general_specialist     25.00s   <- orch_plan_s = 25, round_1 ran 0.00s
```

So the fence a phase happens to carry decided whether it survived, while normal calls in those same turns took 2-13s. Widening fences only converts a failure into a 130s wait; escaping the stall is what recovers the time.

Why screen needs its own fence

Added 2026-08-16, after "stuck at question check" recurred. The call-layer retry above **cannot reach this phase**: its stall fence is 40s and the whole screen budget is 30s, so the phase fence always fires first and the reviewer gets "question check took too long" having had exactly one attempt.

```
the ordering, before
0s      10      20      30      40
|-----|-----|-----|-----|
screen phase budget ..... X
                        ^
                        phase dies here
                        |
                        call-layer stall fence would
                        have fired HERE -- unreachable

after: the fence moves INSIDE the phase
0s      10      30
|-----|-----|
attempt 1 X abandon + re-issue with what is left
      ^
      10s stall fence
```

Three properties, each chosen to avoid re-tuning anything else:

- **The two attempts SHARE the 30s** — the retry gets `30 - elapsed`, so the worst case is unchanged and the existing `screen_timeout` handler still fires on the same `TimeoutError`.
- **It sits at the PHASE, not the call**, so it works on both backends. The call-layer retry is safechain-only; a stall in dev would never have been covered by it.

- **10s is ~2x the phase's own <5s target**, so a healthy screen never trips it. `SCREEN_STALL_RETRY_S=0` disables it. The first attempt is clamped to `min(fence, budget)` — the same bug that had to be fixed in the call layer.

Same tag vocabulary as the call layer, deliberately, so AgenticEval greps one set of names: `stall_retry` / `stall_retry_failed` on the `screen` node, plus `screen_stalled` / `screen_retry_stalled` in the JSONL.

It **retries a stall, it does not fix one**. The re-issue sends the same prompt, so a wedge caused by prompt size will reproduce; `stall_retry_failed` is what distinguishes that from a transient backend hiccup, and `PRIOR_QUESTIONS_FOR_SCREEN` (12) is the knob for the former.

The round-1 watchdog

The oldest of the three, and the one that is easy to mistake for an ordinary phase fence. `orch_plan_s = 25` does NOT bound the planning call — it bounds **time-to-first-tool-call**, and expiry means "this planning attempt is wedged, throw it away and plan again".

```
orchestrator round 1 starts streaming
|
+-- a tool call arrives before 25s -----> PLANNING DONE
|
+-- 25s with no tool call
    | log orchestrator_plan_timeout
    | emit orchestrator_retry (reason: planning_timeout)
    v
    abandon the stream, _orch_attempt += 1, rerun from the top
    |
    +-- THE FINAL ATTEMPT RUNS WITH THE DEADLINE DISARMED
        (_plan_deadline = None when _orch_attempt + 1
         == _MAX_ORCH_ATTEMPTS)
```

The disarm is what makes 25s safe. A tight fence on a single attempt would hard-fail any environment that is merely slow; here the tight fence only ever costs one wasted attempt, and the retry runs unbounded. That is the opposite trade from the reviewer below, which gets one attempt and must therefore be given room rather than speed.

Two other retries share this loop and are worth not confusing with it: `ModelBehaviorError` (the FinalAnswer did not parse) and the dispatch-skip retry (the orchestrator answered without dispatching a specialist), both capped by `_MAX_ORCH_ATTEMPTS = 2`. All three exit to the same `orchestrator_error` / `orchestrator_error_fallback` handler.

How the reviewer runs

The coherence gate on multi-specialist turns. It is not an extra specialist — it reads what the specialists already returned and may send exactly ONE of them back with an anchor.

```
phase 1 done: FinalAnswer + tool_calls
|
+-- fewer than 2 DOMAIN specialists -----> SKIPPED ENTIRELY
    | log general_specialist_skipped          no reviewer, no extra
    |                                         Runner run, zero added
    |                                         latency on the common path
    v
    >= 2 distinct domain specialists
build specialist_outputs from tool_calls
report_agent + general_specialist are EXCLUDED, so the reviewer
never sees its own output and aux tools never reach the gate
|
v
+== FENCED: reviewer_s = 120s =====+
```

```

| Runner.run(general_specialist, {question, outputs}) |
| EVERY exception is swallowed -> None -> review_failed |
| the turn must never block on the reviewer |
+=====+
|
| +-- None (timeout or error) -----> keep the phase-1 answer
|                                     (gate silently absent)
| +-- directive.kind == "coherent" -----> keep the phase-1 answer
|
| +-- "needs_redispatch" + a named specialist
|
|     +-- dispatch_count already 2 -----> keep the phase-1 answer
|     |     log review_capped             + a residual flag saying
|     |                                     it may be unanchored
|     v dispatch_count < 2
|     bump it (the initial dispatch already counted as round 1)
|     discard THAT specialist's phase-1 distillation <- KB hygiene:
|     its mis-anchored KPs have a different topic than the
|     corrected ones, so supersession alone would leave both
|     resume_input = phase-1 transcript + a [REVIEW DIRECTIVE] user turn
|
|     v
+== NOT FENCED =====+
| run_redispatch_pass: Runner.run_streamed on the resume |
| bounded only by specialist_s (240 per agent) and the |
| turn wall clock (360). No planning watchdog either -- |
| that lives in round 1 only. |
+=====+
|
| v
|
| the new FinalAnswer REPLACES the phase-1 one

```

Two things the shape is worth reading for. The fence covers only the reviewer's own call — the re-dispatch it triggers is a whole orchestrator resume running outside it. And the failure mode is silence: a reviewer that times out returns `None`, which is indistinguishable downstream from a reviewer that looked and found the answer coherent. `review_failed` in the JSONL is the only difference.

Why the reviewer needed its own knob

`ORCH_PLAN_TIMEOUT_S` used to fence two things that want opposite treatment.

	on expiry	so it wants
round-1 watchdog	replan; deadline disarmed on the last attempt	25s TIGHT (it retries)
coherence reviewer	caught by a blanket except, <code>review_failed</code> , returns <code>None</code> -- NO retry, the gate is just dropped for the turn	room to OUTLAST a stall (it does not retry)

At 25s the reviewer was also below the 40s call-layer stall fence, so that retry was unreachable for it — the phase fence always won. That is the private- env failure where `general_specialist` died at exactly **25.00s** while `round_1` had run 0.00s.

Split into `reviewer_s = 120`, sized to hold one worst-case call (40s stall fence + 60s retry budget = 100) plus the reviewer's own work. Dev is unaffected: **122 `review_done` events across the shipped logs, zero `review_failed`** — the reviewer never approached 25s there, which is why this only ever showed up in prod.

Measuring this is easy to get wrong: `general_specialist` is ALSO dispatchable as an ordinary specialist under the 240s fence, so a naive duration query pools the two and reports a reviewer p95 of ~26s that does not exist. Filter on the `review` tag, or count `review_done` / `review_failed`.

What a retry costs

An abandoned attempt is not free, and where the cost lands depends on where the wedge is — which is the same unresolved question the call-layer bet rests on.

wedge is...	attempt 1 tokens	what the retry adds
at the provider (inference started)	billed — cancelling closes your side, it does not un-run the work	a second full prompt
before the provider (pool, TLS, socket)	nothing sent, nothing billed	a second full prompt

So worst case is ~2x input and up to 2x output; best case is 1x. For screen the exposure is small (question + <=12 prior questions + the relevance skill), and for a specialist round it is whatever that round's context is.

The trace under-reports this. `attach_usage` records the prompt excerpt *before* the request but the token COUNTS only *after* the response returns (`firewall_client.py:101` vs `:140`). A cancelled call never reaches the second call, so an abandoned attempt contributes **zero tokens to node_trace** while potentially costing real money. AgenticEval scores from the trace, so its token totals are a floor, not a total, on any run where stalls occurred.

Nothing is rebuilt for the retry except the coroutine and the HTTP request. The agent, the shim, `_CLIENTS.firewalled_client`, the httpx pool and (on safechain) the cached LangChain model are all reused — deliberately, since rebuilding a client would add DNS, TLS and token acquisition to a path already in trouble. The consequence: **the retry is fresh at the request level, not below it**. If the wedge lives in the connection or the pool rather than the request, a retry can inherit it.

The multiplication trap

Agent-level and call-level retries compose. A specialist that retries once, each of whose calls stalls once, issues **up to 4 requests** — and its phase fence has to hold all of them. That is why the call-level budget is exactly one extra attempt rather than a general retry loop: the arithmetic has to stay inspectable against the fence above it.

Known inconsistencies

- **RESOLVED** — the "phase fence must be tighter than the per-call cap" invariant was backwards for multi-round agents. It holds for single-call phases (`screen`, `orch_plan`), where a tighter fence gives the clearer error. A specialist making four calls must be LOOSER than one call's cap or it can never complete more than one slow call. The rule is now split by phase kind, and `specialist_s/report_agent_s` are sized to hold one worst-case call ($40 + 140 = 180$) plus the agent's own rounds.
- **RESOLVED** — a phase tighter than the call fence gets no retry at all. The flip side of the above, and it cost real failures before anyone noticed: `screen` at 30s sits under the 40s call-layer stall fence, so that retry was unreachable code for this phase. Fixed by giving the phase its own 10s fence rather than by loosening the phase.

The rule: any single-call phase budgeted under 40s must carry either its own stall fence or a watchdog, or it gets ONE attempt and no retry at all. Nothing tests for this. Three phases sit in that range today and all three are now covered — `screen` (own fence), `orch_plan` (watchdog), and `reviewer`, which was NOT covered until it was moved out of that range entirely.

- **RESOLVED** — the reviewer shared `ORCH_PLAN_TIMEOUT_S` with the round-1 watchdog and the two want opposite things. Split `reviewer_s = 120`; see "Why the reviewer needed its own knob". The claim in the previous edition of this document — that `orch_plan` was "covered by its round-1 watchdog" — was true

of the watchdog and false of the reviewer sharing its knob, which is precisely how the gap survived being written down.

- **The documented "slowest realistic path" does not fit the turn fence, and widening the reviewer made it worse.** $30 + 25 + 240 + 120 + 150 = 565s$ against $\text{turn_wall_clock_s} = 360$ (was 470 when the reviewer was 25). Every term is a worst case that rarely co-occurs, and the turn fence does catch the sum in practice — but the arithmetic is now further from the budget, and the guard test checks each phase against the fence individually, so nothing catches the sum.